

```
## Booting kernel from Legacy Image at a0000000 ...
Image Name: layerscape_kernel
Created: 2017-09-28 9:35:51 UTC
Image Type: AArch64 RTEMS Kernel Image (uncompressed)
Data Size: 42800 Bytes = 41.8 KiB
Load Address: 80010000
Entry Point: 80010000
Verifying Checksum ... OK
Loading Kernel Image ... OK
## Transferring control to RTEMS (at address 80010000) ...
[Kernel] Load module 'ARM-GICv3/4 Wrapper'...Success
[Kernel] Load module 'MMU Controller'...Success
[Kernel] Load module 'Exceptions'...Success
[Kernel] Load module 'Interrupt Handler'...Success
[Kernel] Load module 'Serial UART'...Success
[Kernel] Load module 'Hang'...Success
[Kernel] Load module 'Timer'...Success
[Kernel] Load module 'Syscall Controller'...Success
[Kernel] Load module 'IO Syscalls'...Success
[Kernel] Load module 'Memory Syscalls'...Success
[Kernel] Load module 'Thread Syscalls'...Success
[Kernel] Load module 'Process Scheduler'...Success
[Kernel] Load module 'Multicore-Support'...Success
[Kernel] Load module 'Performance Counter'...Success
[Kernel] Boot completed on all cores...
$:>
```

# Layerscape Kernel

Dominik Walter

University of Augsburg  
Systems and Networking

12.10.2017

# General

- ▶ *freescale LS2085ARDB*
- ▶ 8x ARM Cortex-A57
- ▶ Data path acceleration architecture (DPAA2)
- ▶ 19 GiB RAM
- ▶ 128 MiB NOR Flash
- ▶ 2048 MiB NAND Flash





# Project

Goal: Writing a small baremetal kernel for starting benchmark-applications.

- ▶ C++
- ▶ 3 Code-Projects
  - ▶ layerscape\_kernel
  - ▶ layerscape\_shell
  - ▶ layerscape\_helloworld

# Methode

## How to write a kernel:

- ▶ Step 1: Search for Code-Examples
- ▶ Step 2: Read the Documentation
- ▶ Step 3: "*Try and Error*" until it works
- ▶ Step 4: Repeat

# Installation

- ▶ `install.sh`
  - ▶ Download and install all required binaries.
- ▶ `configure.sh`
  - ▶ Set up the build configuration.

## Required Packages/Binaries:

- ▶ `aarch64-elf/gcc-linaro-5.4.1-2017`
- ▶ `xinetd`
- ▶ `tftpd`
- ▶ `tftp`
- ▶ `u-boot-tools`
- ▶ `make`
- ▶ `cmake`
- ▶ `screen`

# Build Process



# U-Boot

```
U-Boot 2015.10LS2085A-SDK+g3242b20 (Feb 29 2016 - 14:30:27 +0800)

SoC: LS2085E (0x87010010)
Clock Configuration:
    CPU0(A57):1400 MHz  CPU1(A57):1400 MHz  CPU2(A57):1400 MHz
    CPU3(A57):1400 MHz  CPU4(A57):1400 MHz  CPU5(A57):1400 MHz
    CPU6(A57):1400 MHz  CPU7(A57):1400 MHz
    Bus:     400 MHz  DDR:    1866.667 MT/s    DP-DDR:   1600 MT/s
Reset Configuration Word (RCW):
    00: 38303820 38380038 00000000 00000000
    10: 00000000 00200000 00200000 00000000
    20: 01012980 00003200 00000000 00000000
    30: 00000e0b 00000000 00000000 00000000
    40: 00000000 00000000 00000000 00000000
    50: 00000000 00000000 00000000 00000000
    60: 00000000 00000000 00027000 00000000
    70: 412a0000 00000000 00000000 00000000
Model: Freescale Layerscape 2085a RDB Board
Board: LS2085E-RDB, Board Arch: V1, Board version: D, boot from vBank: 4
FPGA: v1.18
SERDES1 Reference : Clock1 = 156.25MHz Clock2 = 156.25MHz
SERDES2 Reference : Clock1 = 100MHz Clock2 = 100MHz
I2C: ready
DRAM: Initializing DDR...using SPD
Detected UDIMM 18ASF1G72AZ-2G3B1
Detected UDIMM 18ASF1G72AZ-2G3B1
```

# Start U-Boot

- ▶ Plug the serial cable into the UART-Port 2
- ▶ Type 'screen /dev/ttyUSB0 115200'
- ▶ Press the Power-Button to boot into U-Boot

```
U-Boot 2015.10.52085A-SDK-g3242b28 (Feb 29 2016 - 14:38:27 +0800)
Set: LS2085E (0x870101010)
Clock Configuration:
CPU(A57):1400 MHz CPU(A57):1400 MHz CPU2(A57):1400 MHz
CPU3(A57):1400 MHz CPU4(A57):1400 MHz CPU5(A57):1400 MHz
CPU6(A57):1400 MHz CPU7(A57):1400 MHz CPU8(A57):1400 MHz
Hus: 400 MHz DR: 1086.667 MHz/s DP-DDR: 1600 MHz/s
Reset Configuration (RCF):
00: 30303020 30303030 00000000 00000000
10: 00000000 00200000 00200000 00000000
20: 00000000 00000000 00000000 00000000
30: 00000000 00000000 00000000 00000000
40: 00000000 00000000 00000000 00000000
50: 00000000 00000000 00000000 00000000
60: 00000000 00000000 00027000 00000000
70: 41240000 00000000 00000000 00000000
Model: LS2085E-SDK Board: LS2085E-SDK Board Arch: V1 Board version: 0, boot from vBank: 4
FPGA: v1.18
EMMC0: Reference : Clock = 156.25MHz Clock2 = 156.25MHz
SERDES2 Reference : Clock = 100MHz Clock2 = 100MHz
TSC: ready
DRAM: initializing 800...using S9P
Detected UDIPM 10A5P1G72A2-2G3B1
Detected UDIPM 10A5P1G72A2-2G3B1
```



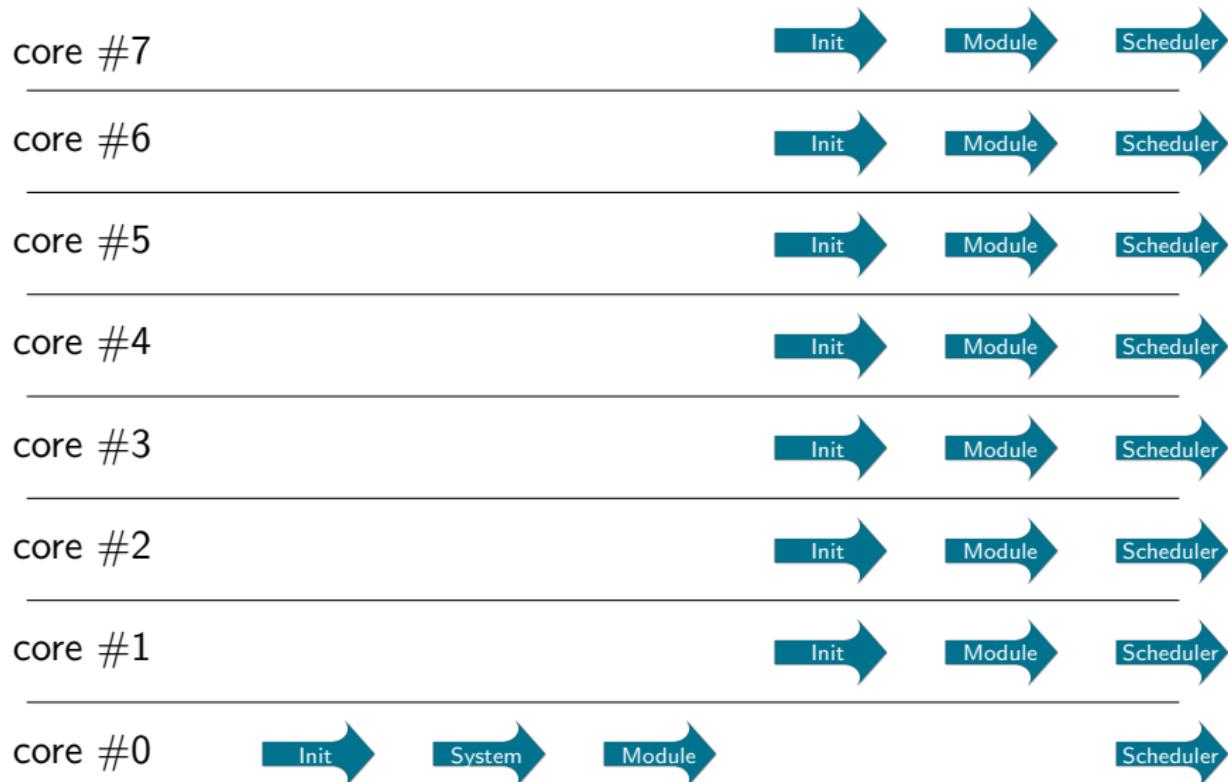
# Deploy images

- ▶ dhcp
- ▶ setenv serverip 137.250.172.39
- ▶ tftpboot file.uimg
- ▶ cp.b \$loadaddr 0xa0000000  
\$filesize
- ▶ bootm 0xa0000000

```
c100000 device
TFTP from server 137.250.16.21; our IP address is 137.250.172.64; sending through gateway 137.250.16.250
Filename: 'pxelinux.0'.
Load address: 0x00100000
Loading...
TFTP error: 'File not found' (1)
Not retrying...
Using c100000 device
TFTP from server 137.250.172.39; our IP address is 137.250.172.64
Filename: 'layerscape kernel.uimg'.
Load address: 0x00100000
Loading: ######/#####
done
Bytes transferred = 42864 (a770 hex)

.dive
Erase 2 sectors
Copy to Flash... 9...8...7...6...5...4...3...2...1...done
Max write speed to Flash...
Un-Protected 1 sectors
Erasing Flash...
done
Erased 1 sectors
Writing to Flash... 9...8...7...69...8...7...6...5...4...3...2...1...done
Protected 1 sectors
done
.no
.tftpboot layerscape shell.img; cp.b 0x00100000 0x00000000 $filesize;
Using c100000 device
TFTP from server 137.250.172.39; our IP address is 137.250.172.64
Filename: 'layerscape shell.img'.
Load address: 0x00100000
Loading: #
done
Bytes transferred = 2688 (a30 hex)
```

# Boot Process



# Applications

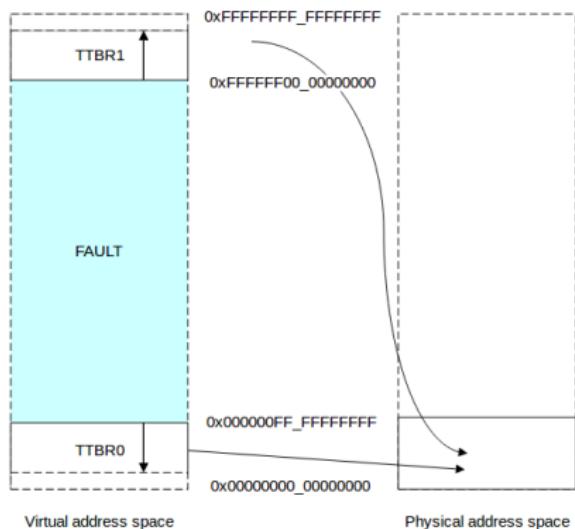
- ▶ Multiple Process-States (e.g. running, blocked)
- ▶ Primitive Signals (e.g. SIGTERM)
- ▶ Various Syscalls
- ▶ Unique Virtual Address Space
- ▶ Multithreading Support
- ▶ No POSIX!

# Start Applications

- ▶ Memory region [0xB0000000, 0xFFFFFFFF] reserved for applications.
- ▶ .img-Format (no ELF-Loader)
- ▶ Load directly from Kernel with `create_process`.
  - ▶ Currently: Image and entrypoint at [0xB0000000] with a maximum size of 4KB.
- ▶ `layerscape_shell` provides a command-line-interface for loading applications.

# Virtual Memory

- ▶ 40 Bit-Address
- ▶ 4KB Granularity
- ▶ Kernel: Leading 0s
- ▶ User: Leading 1s



# Implemented Syscalls

- ▶ Input/Output
  - ▶ `read`
  - ▶ `write`
- ▶ Memory
  - ▶ `malloc`
  - ▶ `free`
- ▶ Process
  - ▶ `pid`
  - ▶ `thread_create`
  - ▶ `thread_join`
  - ▶ `process_create`
  - ▶ `EXIT`

# Outlook

Further Ideas:

- ▶ Power Management
- ▶ Accelerator Unit
- ▶ Network-Stack
- ▶ Filesystem
- ▶ ELF-Loader
- ▶ Scheduling Algorithm
- ▶ Linux Kernelmodule Wrapper
- ▶ POSIX Compatibility